


# Energy-Based Models

Stefano Ermon, Yang Song


Stanford University

Lecture 12

# Recap. of last lecture



$\mathbf{x}_i \sim P_{\text{data}}$   
 $i = 1, 2, \dots, n$



- Energy-based models:  $p_{\theta}(\mathbf{x}) = \frac{\exp\{f_{\theta}(\mathbf{x})\}}{Z(\theta)}$ .
  - $Z(\theta)$  is intractable, so no access to likelihood.
  - Comparing the probability of two points is easy:  
 $p_{\theta}(\mathbf{x}')/p_{\theta}(\mathbf{x}) = \exp(f_{\theta}(\mathbf{x}') - f_{\theta}(\mathbf{x}))$ .
- Maximum likelihood training:  $\max_{\theta} \{f_{\theta}(\mathbf{x}_{\text{train}}) - \log Z(\theta)\}$ .
  - Contrastive divergence:

$$\nabla_{\theta} f_{\theta}(\mathbf{x}_{\text{train}}) - \nabla_{\theta} \log Z(\theta) \approx \nabla_{\theta} f_{\theta}(\mathbf{x}_{\text{train}}) - \nabla_{\theta} f_{\theta}(\mathbf{x}_{\text{sample}}),$$

where  $\mathbf{x}_{\text{sample}} \sim p_{\theta}(\mathbf{x})$ .

# Sampling from EBM: MH-MCMC

Metropolis-Hastings Markov chain Monte Carlo (MCMC).

- ①  $\mathbf{x}^0 \sim \pi(\mathbf{x})$
- ② Repeat for  $t = 0, 1, 2, \dots, T - 1$ :
  - $\mathbf{x}' = \mathbf{x}^t + \text{noise}$
  - $\mathbf{x}^{t+1} = \mathbf{x}'$  if  $f_\theta(\mathbf{x}') \geq f_\theta(\mathbf{x}^t)$
  - If  $f_\theta(\mathbf{x}') < f_\theta(\mathbf{x}^t)$ , set  $\mathbf{x}^{t+1} = \mathbf{x}'$  with probability  $\exp\{f_\theta(\mathbf{x}') - f_\theta(\mathbf{x}^t)\}$ , otherwise set  $\mathbf{x}^{t+1} = \mathbf{x}^t$ .

Properties:

- In theory,  $\mathbf{x}^T$  converges to  $p_\theta(\mathbf{x})$  when  $T \rightarrow \infty$ .
- In practice, need a large number of iterations and convergence slows down exponentially in dimensionality.

# Sampling from EBMs: unadjusted Langevin MCMC

Unadjusted Langevin MCMC:


- ①  $\mathbf{x}^0 \sim \pi(\mathbf{x})$
- ② Repeat for  $t = 0, 1, 2, \dots, T - 1$ :
  - $\mathbf{z}^t \sim \mathcal{N}(0, I)$
  - $\mathbf{x}^{t+1} = \mathbf{x}^t + \epsilon \nabla_{\mathbf{x}} \log p_{\theta}(\mathbf{x})|_{\mathbf{x}=\mathbf{x}^t} + \sqrt{2\epsilon} \mathbf{z}^t$

Properties:


- $\mathbf{x}^T$  converges to  $p_{\theta}(\mathbf{x})$  when  $T \rightarrow \infty$  and  $\epsilon \rightarrow 0$ .
- $\nabla_{\mathbf{x}} \log p_{\theta}(\mathbf{x}) = \nabla_{\mathbf{x}} f_{\theta}(\mathbf{x})$  for continuous energy-based models.
- Convergence slows down as dimensionality grows.

Sampling converges slowly in high dimensional spaces and is thus very expensive, yet we need sampling for **each training iteration** in contrastive divergence.

# Today's lecture



$$\mathbf{x}_i \sim P_{\text{data}} \\ i = 1, 2, \dots, n$$



**Goal:** Training without sampling

- Score Matching
- Noise Contrastive Estimation
- Adversarial training

# Score function

**Energy-based model:**  $p_{\theta}(\mathbf{x}) = \frac{\exp\{f_{\theta}(\mathbf{x})\}}{Z(\theta)}$

**(Stein) Score function:**


$$s_{\theta}(\mathbf{x}) := \nabla_{\mathbf{x}} \log p_{\theta}(\mathbf{x}) = \nabla_{\mathbf{x}} f_{\theta}(\mathbf{x}) - \underbrace{\nabla_{\mathbf{x}} \log Z(\theta)}_{=0} = \nabla_{\mathbf{x}} f_{\theta}(\mathbf{x})$$

- Gaussian distribution

$$p_{\theta}(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$
$$\longrightarrow s_{\theta}(x) = -\frac{x-\mu}{\sigma^2}$$

- Gamma distribution

$$p_{\theta}(x) = \frac{\beta^{\alpha}}{\Gamma(\alpha)} x^{\alpha-1} e^{-\beta x}$$
$$\longrightarrow s_{\theta}(x) = \frac{\alpha-1}{x} - \beta$$



$p_{\theta}(\mathbf{x})$  vs.  $s_{\theta}(\mathbf{x})$

# Score matching

## Observation

$s_\theta(\mathbf{x}) = \nabla_{\mathbf{x}} \log p_\theta(\mathbf{x})$  is independent of the partition function  $Z(\theta)$ .

Fisher divergence between  $p(\mathbf{x})$  and  $q(\mathbf{x})$ :

$$D_F(p, q) := \frac{1}{2} E_{\mathbf{x} \sim p} [\|\nabla_{\mathbf{x}} \log p(\mathbf{x}) - \nabla_{\mathbf{x}} \log q(\mathbf{x})\|_2^2]$$

**Score matching:** minimizing the Fisher divergence between  $p_{\text{data}}(\mathbf{x})$  and the EBM  $p_\theta(\mathbf{x})$

$$\begin{aligned} & \frac{1}{2} E_{\mathbf{x} \sim p_{\text{data}}} [\|\nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x}) - s_\theta(\mathbf{x})\|_2^2] \\ &= \frac{1}{2} E_{\mathbf{x} \sim p_{\text{data}}} [\|\nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x}) - \nabla_{\mathbf{x}} f_\theta(\mathbf{x})\|_2^2] \end{aligned}$$

# Score matching

$$\frac{1}{2} E_{\mathbf{x} \sim p_{\text{data}}} [\|\nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x}) - \nabla_{\mathbf{x}} \log p_{\theta}(\mathbf{x})\|_2^2]$$

How to deal with  $\nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x})$ ? Integration by parts!

$$\begin{aligned} & \frac{1}{2} E_{x \sim p_{\text{data}}} [(\nabla_x \log p_{\text{data}}(x) - \nabla_x \log p_{\theta}(x))^2] \quad (\text{Univariate case}) \\ &= \frac{1}{2} \int p_{\text{data}}(x) [(\nabla_x \log p_{\text{data}}(x) - \nabla_x \log p_{\theta}(x))^2] dx \\ &= \frac{1}{2} \int p_{\text{data}}(x) (\nabla_x \log p_{\text{data}}(x))^2 dx + \frac{1}{2} \int p_{\text{data}}(x) (\nabla_x \log p_{\theta}(x))^2 dx \\ &\quad - \int p_{\text{data}}(x) \nabla_x \log p_{\text{data}}(x) \nabla_x \log p_{\theta}(x) dx \end{aligned}$$

For the cross-correlation term:

$$\begin{aligned} & - \int p_{\text{data}}(x) \nabla_x \log p_{\text{data}}(x) \nabla_x \log p_{\theta}(x) dx \\ &= - \int p_{\text{data}}(x) \frac{1}{p_{\text{data}}(x)} \nabla_x p_{\text{data}}(x) \nabla_x \log p_{\theta}(x) dx \\ &= \underbrace{- p_{\text{data}}(x) \nabla_x \log p_{\theta}(x) \Big|_{x=-\infty}^{\infty}}_{=0} + \int p_{\text{data}}(x) \nabla_x^2 \log p_{\theta}(x) dx \\ &= \int p_{\text{data}}(x) \nabla_x^2 \log p_{\theta}(x) dx \end{aligned}$$



# Score matching

## Univariate score matching

$$\begin{aligned}& \frac{1}{2} E_{\mathbf{x} \sim p_{\text{data}}} [(\nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x}) - \nabla_{\mathbf{x}} \log p_{\theta}(\mathbf{x}))^2] \\&= \frac{1}{2} \int p_{\text{data}}(\mathbf{x}) (\nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x}))^2 d\mathbf{x} + \frac{1}{2} \int p_{\text{data}}(\mathbf{x}) (\nabla_{\mathbf{x}} \log p_{\theta}(\mathbf{x}))^2 d\mathbf{x} \\&\quad - \int p_{\text{data}}(\mathbf{x}) \nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x}) \nabla_{\mathbf{x}} \log p_{\theta}(\mathbf{x}) d\mathbf{x} \\&= \underbrace{\frac{1}{2} \int p_{\text{data}}(\mathbf{x}) (\nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x}))^2 d\mathbf{x}}_{\text{const.}} + \frac{1}{2} \int p_{\text{data}}(\mathbf{x}) (\nabla_{\mathbf{x}} \log p_{\theta}(\mathbf{x}))^2 d\mathbf{x} \\&\quad + \int p_{\text{data}}(\mathbf{x}) \nabla_{\mathbf{x}}^2 \log p_{\theta}(\mathbf{x}) d\mathbf{x} \\&= E_{\mathbf{x} \sim p_{\text{data}}} \left[ \frac{1}{2} (\nabla_{\mathbf{x}} \log p_{\theta}(\mathbf{x}))^2 + \nabla_{\mathbf{x}}^2 \log p_{\theta}(\mathbf{x}) \right] + \text{const.}\end{aligned}$$

## Multivariate score matching

$$\begin{aligned}& \frac{1}{2} E_{\mathbf{x} \sim p_{\text{data}}} [\|\nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x}) - \nabla_{\mathbf{x}} \log p_{\theta}(\mathbf{x})\|_2^2] \\&= E_{\mathbf{x} \sim p_{\text{data}}} \left[ \frac{1}{2} \|\nabla_{\mathbf{x}} \log p_{\theta}(\mathbf{x})\|_2^2 + \text{tr} \left( \underbrace{\nabla_{\mathbf{x}}^2 \log p_{\theta}(\mathbf{x})}_{\text{Hessian of } \log p_{\theta}(\mathbf{x})} \right) \right] + \text{const.}\end{aligned}$$

# Score matching

- 1 Sample a mini-batch of datapoints  $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\} \sim p_{\text{data}}(\mathbf{x})$ .
- 2 Estimate the score matching loss with the empirical mean

$$\begin{aligned} & \frac{1}{n} \sum_{i=1}^n \left[ \frac{1}{2} \|\nabla_{\mathbf{x}} \log p_{\theta}(\mathbf{x}_i)\|_2^2 + \text{tr}(\nabla_{\mathbf{x}}^2 \log p_{\theta}(\mathbf{x}_i)) \right] \\ &= \frac{1}{n} \sum_{i=1}^n \left[ \frac{1}{2} \|\nabla_{\mathbf{x}} f_{\theta}(\mathbf{x}_i)\|_2^2 + \text{tr}(\nabla_{\mathbf{x}}^2 f_{\theta}(\mathbf{x}_i)) \right] \end{aligned}$$

- 3 Stochastic gradient descent.
- 4 No need to sample from the EBM!

## Caveat

Computing the trace of Hessian  $\text{tr}(\nabla_{\mathbf{x}}^2 \log p_{\theta}(\mathbf{x}))$  is in general very expensive for large models.

Denoising score matching (Vincent 2010) and sliced score matching (Song et al. 2019).

# Score matching for learning implicit VAEs


- **Model:**  $p(\mathbf{z}), p_{\theta}(\mathbf{x} | \mathbf{z}), q_{\phi}(\mathbf{z} | \mathbf{x}) = \delta(\mathbf{z} = f_{\phi}(\mathbf{x}, \epsilon))$ .
- **Goal:** maximize the evidence lower bound (ELBO):

$$E_{\mathbf{z} \sim q_{\phi}(\mathbf{z} | \mathbf{x})} [\log p_{\theta}(\mathbf{x} | \mathbf{z}) p(\mathbf{z})] - \underbrace{E_{\mathbf{z} \sim q_{\phi}(\mathbf{z} | \mathbf{x})} \log q_{\phi}(\mathbf{z} | \mathbf{x})}_{:= H(q_{\phi}(\mathbf{z} | \mathbf{x}))}$$

- Estimate the gradient of the entropy term by training an energy-based model.

$$\begin{aligned} & \nabla_{\phi} H(q_{\phi}(\mathbf{z} | \mathbf{x})) \\ &= -\nabla_{\phi} E_{\mathbf{z} \sim q_{\phi}(\mathbf{z} | \mathbf{x})} [\log q_{\phi}(\mathbf{z} | \mathbf{x})] \\ &= -\nabla_{\phi} E_{\epsilon} [\log q_{\phi}(f_{\phi}(\mathbf{x}, \epsilon) | \mathbf{x})] = -E_{\epsilon} [\nabla_{\phi} \log q_{\phi}(f_{\phi}(\mathbf{x}, \epsilon) | \mathbf{x})] \\ &= -E_{\epsilon} [\underbrace{\nabla_{\mathbf{z}} \log q_{\phi}(\mathbf{z} | \mathbf{x})}_{\text{Score function of } q_{\phi}(\mathbf{z} | \mathbf{x})} \big|_{\mathbf{z}=f_{\phi}(\mathbf{x}, \epsilon)} \nabla_{\phi} f_{\phi}(\mathbf{x}, \epsilon)] \end{aligned}$$


# Score matching for learning implicit VAEs




Samples on CelebA  $64 \times 64$ .

Image source: Song et al., 2019.

# Recap.



$$\mathbf{x}_i \sim P_{\text{data}} \\ i = 1, 2, \dots, n$$



Distances used for training energy-based models.

- KL divergence = maximum likelihood.

$$\nabla_{\theta} f_{\theta}(\mathbf{x}_{\text{data}}) - f_{\theta}(\mathbf{x}_{\text{sample}}) \quad (\text{contrastive divergence})$$

- Fisher divergence = score matching.

$$\frac{1}{2} E_{\mathbf{x} \sim p_{\text{data}}} [\|\nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x}) - \nabla_{\mathbf{x}} f_{\theta}(\mathbf{x})\|_2^2]$$

# Noise contrastive estimation

Learning an energy-based model by contrasting it with a noise distribution.

- Data distribution:  $p_{\text{data}}(\mathbf{x})$ .
- Noise distribution:  $p_n(\mathbf{x})$ . Should be analytically tractable and easy to sample from.
- Training a discriminator  $D_\theta(\mathbf{x}) \in [0, 1]$  to distinguish between data samples and noise samples.

$$\max_{\theta} E_{\mathbf{x} \sim p_{\text{data}}} [\log D_\theta(\mathbf{x})] + E_{\mathbf{x} \sim p_n} [\log(1 - D_\theta(\mathbf{x}))]$$

- Optimal discriminator  $D_{\theta^*}(\mathbf{x})$ .

$$D_{\theta^*}(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_n(\mathbf{x})}$$

# Noise contrastive estimation

What if the discriminator is parameterized by

$$D_{\theta}(\mathbf{x}) = \frac{p_{\theta}(\mathbf{x})}{p_{\theta}(\mathbf{x}) + p_n(\mathbf{x})}$$

The optimal discriminator  $D_{\theta^*}(\mathbf{x})$  satisfies

$$D_{\theta^*}(\mathbf{x}) = \frac{p_{\theta^*}(\mathbf{x})}{p_{\theta^*}(\mathbf{x}) + p_n(\mathbf{x})} = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_n(\mathbf{x})}$$

Equivalently,

$$p_{\theta^*}(\mathbf{x}) = \frac{p_n(\mathbf{x}) D_{\theta^*}(\mathbf{x})}{1 - D_{\theta^*}(\mathbf{x})} = p_{\text{data}}(\mathbf{x})$$

# Noise contrastive estimation for training EBMs

Energy-based model:

$$p_{\theta}(\mathbf{x}) = \frac{e^{f_{\theta}(\mathbf{x})}}{Z(\theta)}$$

The constraint  $Z(\theta) = \int e^{f_{\theta}(\mathbf{x})} d\mathbf{x}$  is hard to satisfy.

**Solution:** Modeling  $Z(\theta)$  with an additional trainable parameter  $Z$  that *disregards* the constraint  $Z = \int e^{f_{\theta}(\mathbf{x})} d\mathbf{x}$ .

$$p_{\theta,Z}(\mathbf{x}) = \frac{e^{f_{\theta}(\mathbf{x})}}{Z}$$

With noise contrastive estimation, the optimal parameters  $\theta^*, Z^*$  are

$$p_{\theta^*, Z^*}(\mathbf{x}) = \frac{e^{f_{\theta^*}(\mathbf{x})}}{Z^*} = p_{\text{data}}(\mathbf{x})$$

The optimal parameter  $Z^*$  is the correct partition function, because

$$\int \frac{e^{f_{\theta^*}(\mathbf{x})}}{Z^*} d\mathbf{x} = \int p_{\text{data}}(\mathbf{x}) d\mathbf{x} = 1 \implies Z^* = \int e^{f_{\theta^*}(\mathbf{x})} d\mathbf{x}$$



# Noise contrastive estimation for training EBMs

The discriminator  $D_{\theta,Z}(\mathbf{x})$  for probabilistic model  $p_{\theta,Z}(\mathbf{x})$  is

$$D_{\theta,Z}(\mathbf{x}) = \frac{\frac{e^{f_{\theta}(\mathbf{x})}}{Z}}{\frac{e^{f_{\theta}(\mathbf{x})}}{Z} + p_n(\mathbf{x})} = \frac{e^{f_{\theta}(\mathbf{x})}}{e^{f_{\theta}(\mathbf{x})} + p_n(\mathbf{x})Z}$$

Noise contrastive estimation training

$$\max_{\theta,Z} E_{\mathbf{x} \sim p_{\text{data}}} [\log D_{\theta,Z}(\mathbf{x})] + E_{\mathbf{x} \sim p_n} [\log(1 - D_{\theta,Z}(\mathbf{x}))]$$

Equivalently,

$$\begin{aligned} \max_{\theta,Z} E_{\mathbf{x} \sim p_{\text{data}}} [f_{\theta}(\mathbf{x}) - \log(e^{f_{\theta}(\mathbf{x})} + Zp_n(\mathbf{x}))] \\ + E_{\mathbf{x} \sim p_n} [\log(Zp_n(\mathbf{x})) - \log(e^{f_{\theta}(\mathbf{x})} + Zp_n(\mathbf{x}))] \end{aligned}$$

Log-sum-exp trick for numerical stability:

$$\begin{aligned} \log(e^{f_{\theta}(\mathbf{x})} + Zp_n(\mathbf{x})) &= \log(e^{f_{\theta}(\mathbf{x})} + e^{\log Z + \log p_n(\mathbf{x})}) \\ &= \text{logsumexp}(f_{\theta}(\mathbf{x}), \log Z + \log p_n(\mathbf{x})) \end{aligned}$$

# Noise contrastive estimation for training EBMs

- 1 Sample a mini-batch of datapoints  $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n \sim p_{\text{data}}(\mathbf{x})$ .
- 2 Sample a mini-batch of noise samples  $\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_n \sim p_n(\mathbf{y})$ .
- 3 Estimate the NCE loss.

$$\frac{1}{n} \sum_{i=1}^n [f_{\theta}(\mathbf{x}_i) - \text{logsumexp}(f_{\theta}(\mathbf{x}_i), \log Z + \log p_n(\mathbf{x}_i)) \\ + \log Z + p_n(\mathbf{y}_i) - \text{logsumexp}(f_{\theta}(\mathbf{y}_i), \log Z + \log p_n(\mathbf{y}_i))]$$

- 4 Stochastic gradient ascent.
- 5 No need to sample from the EBM!

# Comparing NCE and GAN

## Similarities:

- Both involve training a discriminator to perform binary classification with a cross-entropy loss.
- Both are likelihood-free.

## Differences:

- GAN requires adversarial training or minimax optimization for training, while NCE does not.
- NCE requires the likelihood of the noise distribution for training, while GAN only requires efficient sampling from the prior.
- NCE trains an energy-based model, while GAN trains a deterministic sample generator.

# Flow contrastive estimation (Gao et al. 2020)

## Observations:

- We need to both evaluate the probability of  $p_n(\mathbf{x})$ , and sample from it efficiently.
- We hope to make the classification task as hard as possible, i.e.,  $p_n(\mathbf{x})$  should be close to  $p_{\text{data}}(\mathbf{x})$  (but not exactly the same).

## Flow contrastive estimation:



- Parameterize the noise distribution with a normalizing flow model  $p_{n,\phi}(\mathbf{x})$ .
- Parameterize the discriminator  $D_{\theta,Z,\phi}(\mathbf{x})$  as

$$D_{\theta,Z,\phi}(\mathbf{x}) = \frac{\frac{e^{f_{\theta}(\mathbf{x})}}{Z}}{\frac{e^{f_{\theta}(\mathbf{x})}}{Z} + p_{n,\phi}(\mathbf{x})} = \frac{e^{f_{\theta}(\mathbf{x})}}{e^{f_{\theta}(\mathbf{x})} + p_{n,\phi}(\mathbf{x})Z}$$

- Train the flow model to minimize  $D_{JS}(p_{\text{data}}, p_{n,\phi})$ :

$$\min_{\phi} \max_{\theta, Z} E_{\mathbf{x} \sim p_{\text{data}}} [\log D_{\theta,Z,\phi}(\mathbf{x})] + E_{\mathbf{x} \sim p_{n,\phi}} [\log(1 - D_{\theta,Z,\phi}(\mathbf{x}))]$$

# Flow contrastive estimation (Gao et al. 2020)



Samples from SVHN, CIFAR-10, and CelebA datasets.

Image source: Gao et al. 2020.

# Adversarial training for EBM

Energy-based model:

$$p_{\theta}(\mathbf{x}) = \frac{e^{f_{\theta}(\mathbf{x})}}{Z(\theta)}$$

Upper bounding its log-likelihood with a variational distribution  $q_{\phi}(\mathbf{x})$ :

$$\begin{aligned} E_{\mathbf{x} \sim p_{\text{data}}}[\log p_{\theta}(\mathbf{x})] &= E_{\mathbf{x} \sim p_{\text{data}}}[f_{\theta}(\mathbf{x})] - \log Z(\theta) \\ &= E_{\mathbf{x} \sim p_{\text{data}}}[f_{\theta}(\mathbf{x})] - \log \int e^{f_{\theta}(\mathbf{x})} d\mathbf{x} \\ &= E_{\mathbf{x} \sim p_{\text{data}}}[f_{\theta}(\mathbf{x})] - \log \int q_{\phi}(\mathbf{x}) \frac{e^{f_{\theta}(\mathbf{x})}}{q_{\phi}(\mathbf{x})} d\mathbf{x} \\ &\leq E_{\mathbf{x} \sim p_{\text{data}}}[f_{\theta}(\mathbf{x})] - \int q_{\phi}(\mathbf{x}) \log \frac{e^{f_{\theta}(\mathbf{x})}}{q_{\phi}(\mathbf{x})} d\mathbf{x} \\ &= E_{\mathbf{x} \sim p_{\text{data}}}[f_{\theta}(\mathbf{x})] - E_{\mathbf{x} \sim q_{\phi}}[f_{\theta}(\mathbf{x})] + H(q_{\phi}(\mathbf{x})) \end{aligned}$$

Adversarial training

$$\max_{\theta} \min_{\phi} E_{\mathbf{x} \sim p_{\text{data}}}[f_{\theta}(\mathbf{x})] - E_{\mathbf{x} \sim q_{\phi}}[f_{\theta}(\mathbf{x})] + H(q_{\phi}(\mathbf{x}))$$

What do we require for the model  $q_{\phi}(\mathbf{x})$ ?

# Conclusion

- Energy-based models are very flexible probabilistic models with intractable partition functions.
- Sampling is hard and typically requires iterative MCMC approaches.
- Computing the likelihood is hard.
- Comparing the likelihood/probability of two different points is tractable.
- Maximum likelihood training by contrastive divergence. Requires sampling for each training iteration.
- Sampling-free training: score matching.
- Sampling-free training: noise contrastive estimation. Additionally provides an estimate of the partition function.
- Sampling-free training: adversarial optimization.
- Reference: *How to Train Your Energy-Based Models* by Yang Song and Durk Kingma.